

Multi-Headed Attention (MHA)



This is a tutorial/implementation of multi-headed attention from paper [Attention Is All You Need](#) in [PyTorch](#). The implementation is inspired from [Annotated Transformer](#).

Here is the [training code](#) that uses a basic transformer with MHA for NLP auto-regression.

[Here is an experiment implementation](#) that trains a simple transformer.

```
24 import math
25 from typing import Optional, List
26
27 import torch
28 from torch import nn
29
30 from labml import tracker
```

Prepare for multi-head attention

This module does a linear transformation and splits the vector into given number of heads for multi-head attention. This is used to transform **key**, **query**, and **value** vectors.

```
33 class PrepareForMultiHeadAttention(nn.Module):
```

	<pre> 44 def __init__(self, d_model: int, heads: int, d_k: int, bias: bool): 45 super().__init__() </pre>
Linear layer for linear transform	<pre> 47 self.linear = nn.Linear(d_model, heads * d_k, bias=bias) </pre>
# Number of heads	<pre> 49 self.heads = heads </pre>
Number of dimensions in vectors in each head	<pre> 51 self.d_k = d_k </pre>
	<pre> 53 def forward(self, x: torch.Tensor): </pre>
Input has shape <code>[seq_len, batch_size, d_model]</code> or <code>[batch_size, d_model]</code> . We apply the linear transformation to the last dimension and split that into the heads.	<pre> 57 head_shape = x.shape[:-1] </pre>
Linear transform	<pre> 60 x = self.linear(x) </pre>
Split last dimension into heads	<pre> 63 x = x.view(*head_shape, self.heads, self.d_k) </pre>
Output has shape <code>[seq_len, batch_size, heads, d_k]</code> or <code>[batch_size, heads, d_model]</code>	<pre> 66 return x </pre>

Multi-Head Attention Module

This computes scaled multi-headed attention for given `query`, `key` and `value` vectors.

```

69 class MultiHeadAttention(nn.Module):

```

$$Attention(Q, K, V) = \underset{seq}{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

In simple terms, it finds keys that matches the query, and gets the values of those keys.

It uses dot-product of query and key as the indicator of how matching they are. Before taking the *softmax* the dot-products are scaled by $\frac{1}{\sqrt{d_k}}$. This is done to avoid large dot-product values causing softmax to give very small gradients when d_k is large.

Softmax is calculated along the axis of of the sequence (or time).

- `heads` is the number of heads.
- `d_model` is the number of features in the `query`, `key` and `value` vectors.

```
90     def __init__(self, heads: int, d_model: int, dropout_prob: float = 0.1, bias: bool
= True):
```

```
96         super().__init__()
```

Number of features per head

```
99         self.d_k = d_model // heads
```

Number of heads

```
101        self.heads = heads
```

These transform the `query`, `key` and `value` vectors for multi-headed attention.

```
104        self.query = PrepareForMultiHeadAttention(d_model, heads, self.d_k, bias=bias)
105        self.key = PrepareForMultiHeadAttention(d_model, heads, self.d_k, bias=bias)
106        self.value = PrepareForMultiHeadAttention(d_model, heads, self.d_k, bias=True)
```

Softmax for attention along the time dimension of <code>key</code>	109	<code>self.softmax = nn.Softmax(dim=1)</code>
Output layer	112	<code>self.output = nn.Linear(d_model, d_model)</code>
Dropout	114	<code>self.dropout = nn.Dropout(dropout_prob)</code>
Scaling factor before the softmax	116	<code>self.scale = 1 / math.sqrt(self.d_k)</code>
We store attentions so that it can be used for logging, or other computations if needed	119	<code>self.attn = None</code>
Calculate scores between queries and keys This method can be overridden for other variations like relative attention.	121	<code>def get_scores(self, query: torch.Tensor, key: torch.Tensor):</code>
Calculate QK^T or $S_{ijbh} = \sum_d Q_{ibhd} K_{jbhd}$	129	<code>return torch.einsum('ibhd,jbhd->ijbh', query, key)</code>
<code>mask</code> has shape <code>[seq_len_q, seq_len_k, batch_size]</code> , where first dimension is the query dimension. If the query dimension is equal to 1 it will be broadcasted.	131	<code>def prepare_mask(self, mask: torch.Tensor, query_shape: List[int], key_shape: List[int]):</code>
	137	<code>assert mask.shape[0] == 1 or mask.shape[0] == query_shape[0]</code>
	138	<code>assert mask.shape[1] == key_shape[0]</code>
	139	<code>assert mask.shape[2] == 1 or mask.shape[2] == query_shape[1]</code>

Same mask applied to all heads.

142

```
mask = mask.unsqueeze(-1)
```

resulting mask has shape

145

```
return mask
```

```
[seq_len_q, seq_len_k, batch_size, heads]
```

`query`, `key` and `value` are the tensors that store collection of *query*, *key* and *value* vectors. They have shape `[seq_len, batch_size, d_model]`.

147

```
def forward(self, *,
```

148

```
    query: torch.Tensor,
```

149

```
    key: torch.Tensor,
```

150

```
    value: torch.Tensor,
```

151

```
    mask: Optional[torch.Tensor] = None):
```

`mask` has shape `[seq_len, seq_len, batch_size]` and

`mask[i, j, b]` indicates whether for batch `b`, query at position `i` has access to key-value at position `j`.

`query`, `key` and `value` have shape

163

```
seq_len, batch_size, _ = query.shape
```

```
[seq_len, batch_size, d_model]
```

164

```
if mask is not None:
```

165

```
    mask = self.prepare_mask(mask, query.shape, key.shape)
```

166

Prepare `query`, `key` and `value` for attention computation. These will then have shape

170

```
query = self.query(query)
```

171

```
key = self.key(key)
```

172

```
value = self.value(value)
```

```
[seq_len, batch_size, heads, d_k]
```

Compute attention scores QK^T . This gives a tensor of shape

176

```
scores = self.get_scores(query, key)
```

```
[seq_len, seq_len, batch_size, heads]
```

Scale scores $\frac{QK^T}{\sqrt{d_k}}$

179

```
scores *= self.scale
```

Apply mask

182

```
if mask is not None:
```

183

```
    scores = scores.masked_fill(mask == 0, float('-inf'))
```

